

# Trees



Statistical Learning  
Master in Big Data. University of Santiago de Compostela  
Manuel Mucientes

# Background

- Stratify or segment the input space into a number of simple regions
- Splitting rules summarized in a decision tree
- Predictions of new observations: mean or mode of the training observations in the region
- Simple and useful for interpretation
- Not competitive with best supervised learning approaches
- Dramatic improvements when used in combination with bagging or boosting
- We focus the discussion on CART: classification and regression tree

# Example

- Predicting the baseball player's salaries using regression trees
  - Inputs: years, hits, etc.
  - Output: salary (log-transformed)
    - $R_1$  mean log salary (\$1,000) of \$165,174 (5.107);  $R_2$  \$402,834;  $R_3$  \$845,346
  - Terminal nodes or leaves, internal nodes, branches
  - Easy to interpret, nice graphical representation



# Regression Trees

- Two steps:

- Divide the predictor space (set of possible values of  $X_1, \dots, X_p$ ) into  $J$  distinct and non-overlapping regions,  $R_1, \dots, R_J$
- For each observation that falls in  $R_j$ , the prediction will be the mean of the output values of the training observations in  $R_j$

- How do we construct the regions,  $R_1, \dots, R_J$ ?

- Regions are boxes (high dimensional rectangles)
- Find boxes that minimize RSS:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

- Computationally infeasible to consider every partition into  $J$  boxes

# Regression Trees (ii)





- Top-down greedy approach: recursive binary splitting
  - For each box, select the predictor  $X_j$  and the cutpoint  $s$  to minimize RSS within each region:

$$R_1(j, s) = \{X | X_j < s\} \quad \text{and} \quad R_2(j, s) = \{X | X_j \geq s\}$$

$$\sum_{i: x_i \in R_1(j, s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j, s)} (y_i - \hat{y}_{R_2})^2$$

- Continue until stopping criterion is reached: usually a minimum node size
- Predict the response of a new test observation using the mean of the training observations in the region
  - Predict the confidence: standard deviation

# Example



# Tree Pruning


- The previous method generates complex trees: overfitting
- Smaller tree: lower variance, better interpretation, slightly higher bias
- One possible solution: build the tree so long as the decrease in RSS due to a split exceeds some threshold
  - Smaller trees, but short sighted strategy: a worthless split may be followed by a very good split later on
- A better strategy: grow a large tree, and prune it back
  - Best way to prune?
    - Cross-validation: extremely large number of subtrees
    - Cost complexity pruning (weakest link pruning)

# Cost Complexity Pruning

- Collapse the internal node that produces the smallest per-node increase in:

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2$$



- Sequence of trees with different sizes
- Example: hitters data
  - Tree learned with 9 features
  - Training: 132 examples
  - Test: 131 examples
- Unpruned tree





# Example (hitters data)

- Six-fold cross-validation (number of examples multiple of 6)
- Minimum error: three node tree



# Classification Trees

- Similar to a regression tree, but with qualitative response
- Predicted class: most commonly occurring class of training observations in the region to which it belongs
  - We are also interested in the class proportions among the training observations in that region
- Tree growing: recursive binary splitting (as in regression)
  - RSS cannot be used
  - A natural alternative: classification error rate
    - Fraction of the training observations of a region that do not belong to the most common class
    - $E = 1 - \max_k(\hat{p}_{mk})$  for the m-th region
    - Not sufficiently sensitive for tree-growing

# Classification Trees (ii)

- Two other measures are preferable

- Gini index:

- $$G = \sum_{k=1}^K \hat{p}_{mk}(1 - \hat{p}_{mk})$$


- Total variance across the K classes

- Cross-entropy:

- $$D = - \sum_{k=1}^K \hat{p}_{mk} \log \hat{p}_{mk}$$

*Node impurity measures for two-class classification*

*Cross-entropy has been scaled*

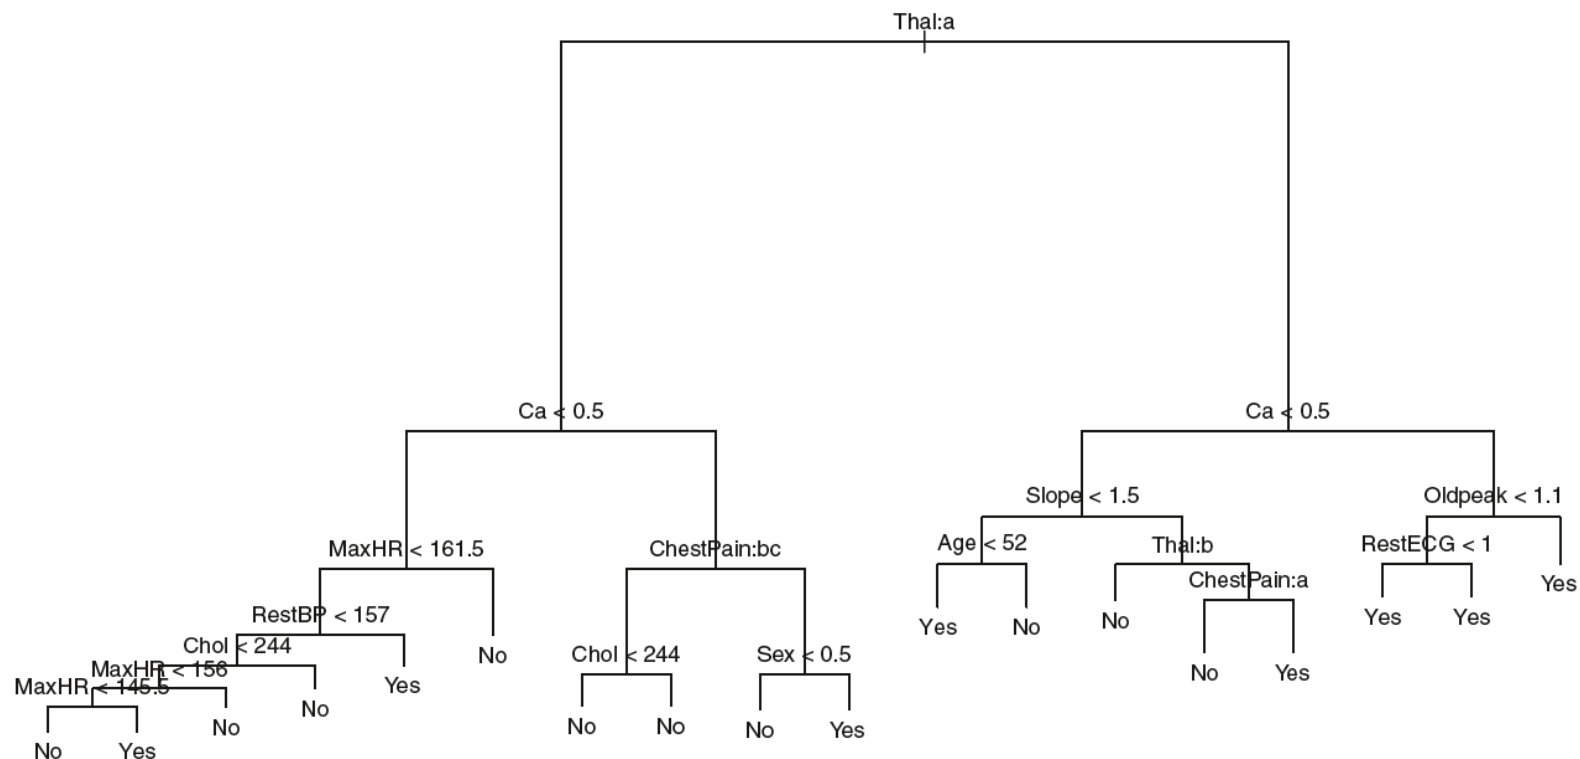


# Classification Trees (iii)

- We need to weight the node impurity measures by the number of observations in the two created child nodes
- For building the tree:
  - Gini index and cross-entropy are preferable: more sensitive to node purity
- For pruning the tree:
  - Any of the three approaches might be used
  - Classification error rate is preferable if prediction accuracy of the final pruned tree is the goal

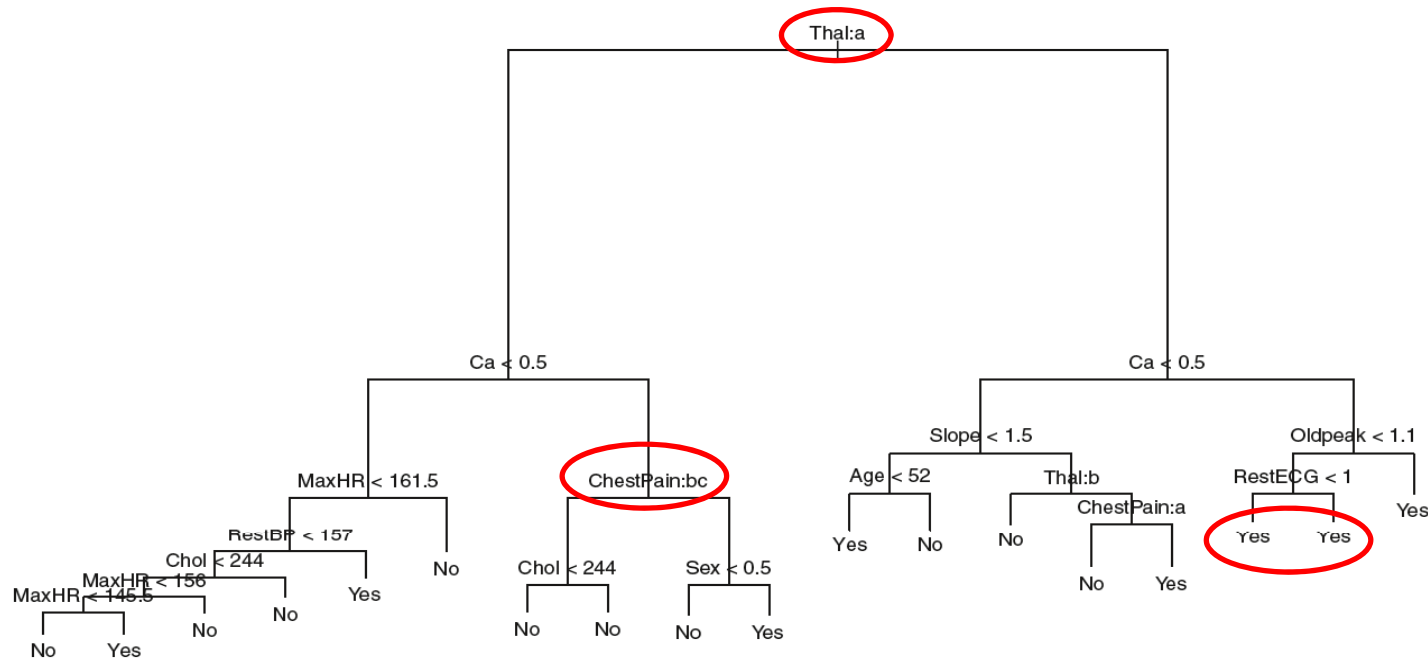
# Example: Heart dataset

- Binary outcome (HD) for 303 patients
  - HD: Yes (heart disease) or No
- 13 predictors: Age, Sex, Chol (cholesterol measurement), Thal (Thalium stress test), ChestPain, ...
  - Categorical (qualitative) predictors: Sex, Thal, ChestPain



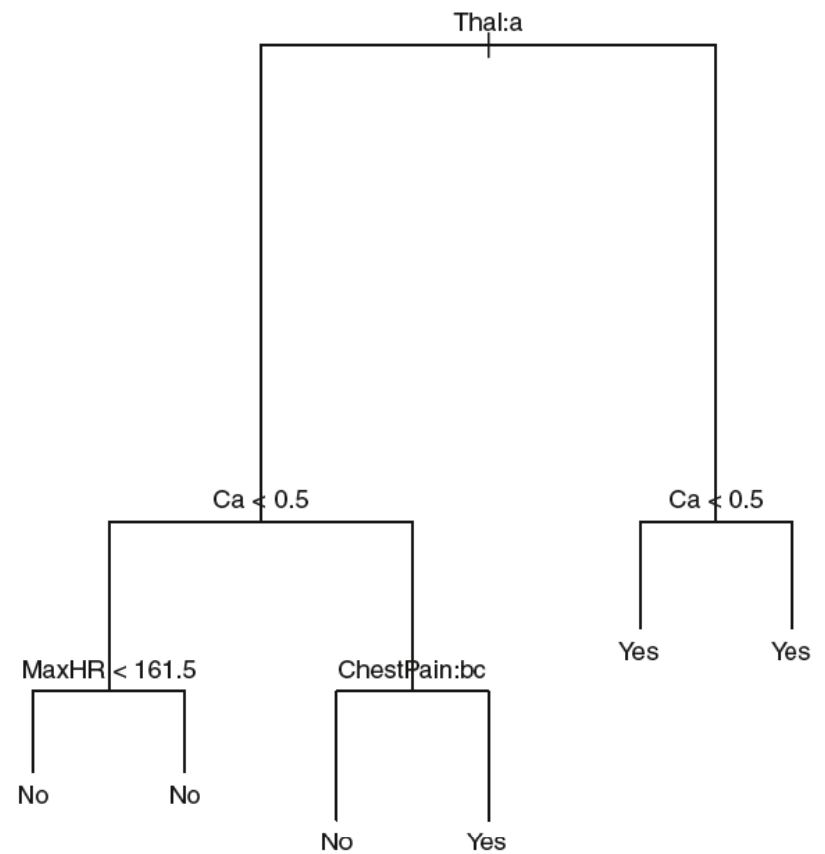

# Example: Heart dataset (ii)

- ChestPain: (a) typical angina, (b) atypical angina, (c) non-anginal pain, (d) asymptomatic
- RestECG: increased node purity
  - Improves Gini index and cross-entropy
  - Classification error not improved
  - Right-hand leaf: 9/9 observations with response value Yes
  - Left-hand leaf: 7/11 observations with response value Yes



# Example: Heart dataset (iii)

- Best tree: six nodes



# Categorical predictors


- $q$  values:  $2^{q-1}-1$  possible partitions into two groups
- In a two-class problem: order the predictor classes according to the proportion falling in outcome class 1
  - Then split as an ordered predictor
  - Optimal split in terms of Gini index and cross-entropy
- This also holds for regression (RSS)
  - Order the categories by increasing mean of the outcome
- For multi-category outcomes no such simplifications are possible
- Try to avoid variables with large  $q$ : overfitting




# Trees vs. Linear Models

- Linear regression:  $f(X) = \beta_0 + \sum_{j=1}^p X_j \beta_j$
- Regression trees:  $f(X) = \sum_{m=1}^M c_m \cdot 1_{(X \in R_m)}$
- A two-dimensional classification example:

*Linear model is better*



*Tree is better*



# Advantages and Disadvantages of trees

- Pros:
  - Easy to explain
  - More closely mirror human decision-making
  - Can be displayed graphically, and are easily interpreted
  - Can handle qualitative predictors without the need to create dummy variables
  
- Cons:
  - Predictive accuracy is lower than other approaches

# Bibliography

- G. James, D. Witten, T. Hastie, y R. Tibshirani, An Introduction to Statistical Learning with Applications in R. Springer, 2013.
  - Chapter 8, Sec. 8.1.
  
- T. Hastie, R. Tibshirani, y J. Friedman, The elements of statistical learning. Springer, 2009.
  - Chapter 9, Sec. 9.2.