

# Neural Networks


Statistical Learning

Master in Big Data. University of Santiago de Compostela

Manuel Mucientes

# Introduction

- Basic idea:
  - Extract linear combinations of the inputs as derived features
  - Model the target as a nonlinear function of these features
- An example of a single neuron:



- $$h_{W,b}(x) = f(W^T x) = f(\sum_{i=1}^3 W_i x_i + b)$$
  - $f$  is the activation function
  - $b$  is the bias term


# Activation functions

- Sigmoid:  $[0, 1]$

$$f(z) = \frac{1}{1 + \exp(-z)}$$

- Derivative:

- $f'(z) = f(z)(1 - f(z))$




# Activation functions (ii)

- Hyperbolic tangent:  $[-1, 1]$

$$f(z) = \tanh(z) = \frac{e^z - e^{-z}}{e^z + e^{-z}}$$

- Derivative:

- $f'(z) = 1 - (f(z))^2$



# Activation functions (iii)

## ■ Rectified linear unit (ReLU):

$$f(z) = \max(0, z)$$


### ■ Derivative:

- 0 if  $z < 0$ , 1 otherwise
- Undefined at  $z = 0$
- Average the gradient over many training examples during optimization

## ■ Gaussian radial basis functions


### ■ RBF networks

$$f(z) = \exp\left(-\frac{1}{2}(z-c)^T \Sigma^{-1}(z-c)\right)$$




# Neural Network Model

- A single hidden layer feed-forward neural network:




- Input layer, hidden layer, output layer
- Parameters:  $(W, b) = (W^{(1)}, b^{(1)}, W^{(2)}, b^{(2)})$ 
  - $W_{ij}^{(l)}$ : weight associated with the connection between unit  $j$  in layer  $l$  and unit  $i$  in layer  $l+1$ 
    - $W^{(1)} \in \mathbb{R}^{3 \times 3}$ , and  $W^{(2)} \in \mathbb{R}^{1 \times 3}$
  - $b_i^{(l)}$  is the bias associated with unit  $i$  in layer  $l+1$ 
    - Bias units do not have inputs or connections going into them

# Neural Network Model (ii)



- $a_i^{(l)}$ : activation (output value) of unit  $i$  in layer  $l$ 
  - $a_i^{(1)} = x_i$
- $s_l$ : number of nodes in layer  $l$  (not counting the bias unit)

# Neural Network Model (iii)




## ■ Computation:

$$\begin{aligned}a_1^{(2)} &= f(W_{11}^{(1)}x_1 + W_{12}^{(1)}x_2 + W_{13}^{(1)}x_3 + b_1^{(1)}) \\a_2^{(2)} &= f(W_{21}^{(1)}x_1 + W_{22}^{(1)}x_2 + W_{23}^{(1)}x_3 + b_2^{(1)}) \\a_3^{(2)} &= f(W_{31}^{(1)}x_1 + W_{32}^{(1)}x_2 + W_{33}^{(1)}x_3 + b_3^{(1)}) \\h_{W,b}(x) &= a_1^{(3)} = f(W_{11}^{(2)}a_1^{(2)} + W_{12}^{(2)}a_2^{(2)} + W_{13}^{(2)}a_3^{(2)} + b_1^{(2)})\end{aligned}$$



# Neural Network Model (iv)




- If  $z_i^{(l+1)}$  is the total weighted sum of inputs to unit  $i$  in layer  $l+1$ :

$$z_i^{(l+1)} = \sum_{j=1}^{s_l} W_{ij}^{(l)} a_j^{(l)} + b_i^{(l)}$$

- $a_i^{(l)} = f(z_i^{(l)})$

# Neural Network Model (v)



- Computation: forward propagation


$$\begin{aligned}z^{(2)} &= W^{(1)}x + b^{(1)} \\a^{(2)} &= f(z^{(2)}) \\z^{(3)} &= W^{(2)}a^{(2)} + b^{(2)} \\h_{W,b}(x) &= a^{(3)} = f(z^{(3)})\end{aligned}$$

- In general (matrix-vector operations):

$$\begin{aligned}z^{(l+1)} &= W^{(l)}a^{(l)} + b^{(l)} \\a^{(l+1)} &= f(z^{(l+1)})\end{aligned}$$

# Architectures

- Patterns of connectivity between neurons
  - Multiple hidden layers
  - Fully (densely) vs. locally connected
  - Weight sharing (locally connected)
  - feed-forward (no loops)
- Most common choice: multilayer feed-forward neural network (multilayer perceptron network)
  - Forward propagation step to calculate outputs of each layer



# Architectures (ii)

- Regression:
  - Typically, one output unit, except when there are several outputs
  - The output function is typically the identity function
- Classification:
  - For K-class classification, K units in the output layer

E.g.  $\begin{bmatrix} 1 \\ 0 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 1 \\ 0 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 1 \\ 0 \end{bmatrix}, \begin{bmatrix} 0 \\ 0 \\ 0 \\ 1 \end{bmatrix}$   
pedestrian car motorcycle truck

- Output function: *softmax*

$$f(z_i^{(l)}) = \frac{e^{z_i^{(l)}}}{\sum_{j=1}^K e^{z_j^{(l)}}}$$

# Architectures (iii)

- An example of the output of *softmax*
  - Three categories (K=3): bike, car, truck

$$f(z_i^{(l)}) = \frac{e^{z_i^{(l)}}}{\sum_{j=1}^K e^{z_j^{(l)}}}$$



	$z_i$	$\exp(z_i)$	$f(z_i)$	Correct probs.
Bike	-0.2	0.8	0.02 (2%)	0.00
Car	<b>3.6</b>	<b>36.6</b>	<b>0.75 (75%)</b>	<b>1.00</b>
Truck	2.4	11.0	0.23 (23%)	0.00

# Fitting Neural Networks

- Backpropagation algorithm

- Training examples:  $\{(x^{(1)}, y^{(1)}), \dots, (x^{(m)}, y^{(m)})\}$

- For a single example:

- Regression and classification:

$$J(W, b; x, y) = \frac{1}{2} \|h_{W,b}(x) - y\|^2$$

- For classification (mutually exclusive classes), cross-entropy is also valid:

$$J(W, b, x, y) = - \sum_{k=1}^K y_k \log (h_{W,b}(x))_k$$

- For classification:

- Sigmoid, *softmax*:  $y=0, 1$ ; output in  $[0, 1]$

- tanh:  $y=-1, 1$ ; output in  $[-1, 1]$

- For regression: scale the outputs (range depends on the activation function)

# Fitting Neural Networks (ii)

- Overall cost function:

$$\begin{aligned} J(W, b) &= \left[ \frac{1}{m} \sum_{i=1}^m J(W, b; x^{(i)}, y^{(i)}) \right] + \frac{\lambda}{2} \sum_{l=1}^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \left( W_{ji}^{(l)} \right)^2 \\ &= \left[ \frac{1}{m} \sum_{i=1}^m \left( \frac{1}{2} \| h_{W,b}(x^{(i)}) - y^{(i)} \|^2 \right) \right] + \frac{\lambda}{2} \sum_{l=1}^{n_l-1} \sum_{i=1}^{s_l} \sum_{j=1}^{s_{l+1}} \left( W_{ji}^{(l)} \right)^2 \end{aligned}$$

- Regularization term (weight decay): prevent overfitting
  - Not applied to the bias terms
- Minimize  $J$  through backpropagation
  - $J$  is non-convex: local minima
  - Gradient descent: usually works fairly well

# Fitting Neural Networks (iii)

- One iteration of gradient descent updates  $W, b$ : Widrow-Hoff rule

$$W_{ij}^{(l)} = W_{ij}^{(l)} - \alpha \frac{\partial}{\partial W_{ij}^{(l)}} J(W, b)$$
$$b_i^{(l)} = b_i^{(l)} - \alpha \frac{\partial}{\partial b_i^{(l)}} J(W, b)$$

- $\alpha$ : learning rate

- Compute the partial derivatives: backpropagation algorithm

$$\frac{\partial J(W, b)}{\partial W_{ij}^{(l)}} = \left[ \frac{1}{m} \sum_{e=1}^m \frac{\partial J(W, b; x^{(e)}, y^{(e)})}{\partial W_{ij}^{(l)}} \right] + \lambda W_{ij}^{(l)}$$

$$\frac{\partial J(W, b)}{\partial b_i^{(l)}} = \frac{1}{m} \sum_{e=1}^m \frac{\partial J(W, b; x^{(e)}, y^{(e)})}{\partial b_i^{(l)}}$$



# Backpropagation algorithm

1. Perform a feedforward pass, computing the activations for layers  $L_2, L_3$ , and so on up to the output layer  $L_{n_l}$ .

2. For each output unit  $i$  in layer  $n_l$  (the output layer), set

$$\delta_i^{(n_l)} = \frac{\partial}{\partial z_i^{(n_l)}} \frac{1}{2} \|y - h_{W,b}(x)\|^2 = -(y_i - a_i^{(n_l)}) \cdot f'(z_i^{(n_l)})$$

3. For  $l = n_l - 1, n_l - 2, n_l - 3, \dots, 2$

For each node  $i$  in layer  $l$ , set

$$\delta_i^{(l)} = \left( \sum_{j=1}^{s_{l+1}} W_{ji}^{(l)} \delta_j^{(l+1)} \right) f'(z_i^{(l)})$$

Intuition:

$\delta_i^{(l)}$  = error of node  $i$  in layer  $l$

4. Compute the desired partial derivatives, which are given as:

$$\begin{aligned} \frac{\partial}{\partial W_{ij}^{(l)}} J(W, b; x, y) &= a_j^{(l)} \delta_i^{(l+1)} \\ \frac{\partial}{\partial b_i^{(l)}} J(W, b; x, y) &= \delta_i^{(l+1)}. \end{aligned}$$

How much error changes with modifications in  $W_{ij}^{(l)}$ :  
[activation of node  $j$  in layer  $l$ ] times  
[error of node  $i$  in layer  $l+1$ ]

# Backpropagation algorithm (ii)

## ■ Matrix-vectorial notation

1. Perform a feedforward pass, computing the activations for layers  $L_2, L_3$ , up to the output layer  $L_{n_l}$ , using the equations defining the forward propagation steps

2. For the output layer (layer  $n_l$ ), set

$$\delta^{(n_l)} = -(y - a^{(n_l)}) \bullet f'(z^{(n_l)})$$

3. For  $l = n_l - 1, n_l - 2, n_l - 3, \dots, 2$ , set

$$\delta^{(l)} = \left( (W^{(l)})^T \delta^{(l+1)} \right) \bullet f'(z^{(l)})$$

4. Compute the desired partial derivatives:

$$\begin{aligned}\nabla_{W^{(l)}} J(W, b; x, y) &= \delta^{(l+1)} (a^{(l)})^T, \\ \nabla_{b^{(l)}} J(W, b; x, y) &= \delta^{(l+1)}.\end{aligned}$$

# Full gradient descent algorithm

1. Set  $\Delta W^{(l)} := 0$ ,  $\Delta b^{(l)} := 0$  (matrix/vector of zeros) for all  $l$ .

2. For  $i = 1$  to  $m$ ,

1. Use backpropagation to compute

$$\nabla_{W^{(l)}} J(W, b; x, y) \text{ and } \nabla_{b^{(l)}} J(W, b; x, y).$$

2. Set  $\Delta W^{(l)} := \Delta W^{(l)} + \nabla_{W^{(l)}} J(W, b; x, y)$ .

3. Set  $\Delta b^{(l)} := \Delta b^{(l)} + \nabla_{b^{(l)}} J(W, b; x, y)$ .

3. Update the parameters:

$$W^{(l)} = W^{(l)} - \alpha \left[ \left( \frac{1}{m} \Delta W^{(l)} \right) + \lambda W^{(l)} \right]$$

$$b^{(l)} = b^{(l)} - \alpha \left[ \frac{1}{m} \Delta b^{(l)} \right]$$


# Training Neural Networks

- Over-parametrized model, non-convex and unstable optimization problem
- Starting values:
  - Random values near 0 for W, biases to 0
    - With standardized inputs, random uniform weights in  $[-0.7, 0.7]$ 
      - Ok for small networks
    - Xavier initialization:  $\text{random()} * \sqrt{1/n}$ 
      - $\text{random()}$ : mean 0, variance 1
      - $n$ : number of inputs
  - Model starts out nearly linear, and becomes nonlinear as weights increase
  - 0 weights give perfect symmetry, large weights lead poor solutions


# Training Neural Networks (ii)

- Overfitting: regularization term (weight decay)
  - Cross-validation to estimate the regularization parameter

Neural Network - 10 Units, No Weight Decay



Neural Network - 10 Units, Weight Decay=0.02



# Training Neural Networks (iii)

## ■ Scaling of the inputs:

- Can have a large effect in the quality of the final solution: scaling of the weights
- Standardize inputs (and outputs) to have mean zero and standard deviation one:
  - Treat all inputs equally: regularization
  - Choose a meaningful range for the starting weights

## ■ Number of hidden units:


- Better to have too many hidden units than too few
  - Too few: not enough flexibility to capture nonlinearities
  - Too many: extra weights can be shrunk toward zero with regularization
- Typical number of hidden units: 5 to 100 (per layer)
  - Number increasing with number of inputs and number of training cases
  - Same number in every layer (reasonable default)

# Training Neural Networks (iv)

- Number of hidden layers:
  - background knowledge and experimentation
  - Allow the construction of hierarchical features at different levels of resolution
- Multiple minima: non-convex error function
  - Try a number of random starting configurations: choose the best solution
  - Better approach: average the predictions over a collection of networks
  - Another approach: bagging
- Batch learning: all the examples
  - Stochastic Gradient Descent (SGD): 1 example
- Advanced optimization methods:
  - Conjugate gradient
  - Quasi-Newton: Broyden–Fletcher–Goldfard–Shanno (BFGS) and L-BFGS

# Example: ZIP Code Data


- Le Cun, 1989
- 16x16 grayscale images
- Training with 320 digits, test with 160
- Five different networks: sigmoidal output units, sum-of-squares error function






# Example: ZIP Code Data (ii)

- Net-1: no hidden layer, equivalent to multinomial logistic regression
- Net-2: one hidden layer, 12 hidden units fully connected
- Net-3: two hidden layers locally connected
  - First hidden layer:
    - Inputs from a 3x3 patch
    - Units 1 unit apart are 2 pixels apart
  - Second hidden layer:
    - Inputs from a 5x5 patch
    - Units 1 unit apart are 2 pixels apart
  - Local connectivity:
    - Each unit extracts local features from the previous layer
    - Lower number of weights




# Example: ZIP Code Data (iii)

- Net-4: two hidden layers, locally connected, weight sharing
  - First hidden layer: two 8x8 feature maps
    - Input from 3x3 patch
    - Units in the same 8x8 feature map share the same set of 9 weights (bias not shared)
    - Extracted features in different parts of the image are computed with the same linear functional:  
**convolutional networks**
  - Second hidden layer: no weight sharing
- The gradient of the error for a shared weight is the sum of the gradients to each connection controlled by the weight



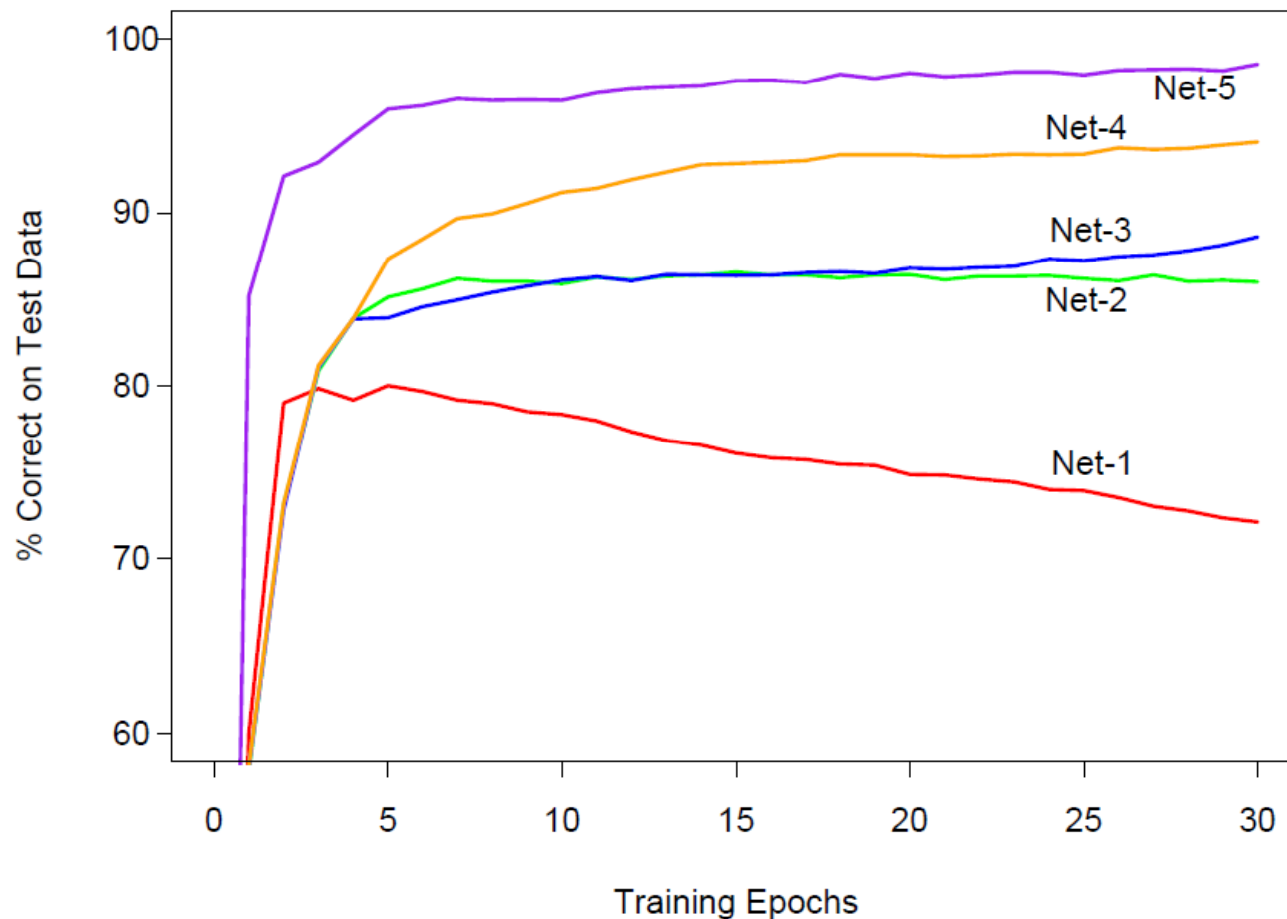
# Example: ZIP Code Data (iv)

- Net-5: two hidden layers, locally connected, two levels of weight sharing
  - First hidden layer: same as Net-4
  - Second hidden layer: four 4x4 feature maps
    - Input from a 5x5 patch
    - Weights shared in each of the feature maps
- Features of handwritten style should appear in more than one part of a digit
- Subject matter knowledge should be used to improve performance



# Example: ZIP Code Data (v)

- Training error was 0% in all the cases
  - More parameters than training observations



# Example: ZIP Code Data (vi)

	Network Architecture	Links	Weights	% Correct
Net-1:	Single layer network	2570	2570	80.0%
Net-2:	Two layer network	3214	3214	87.0%
Net-3:	Locally connected	1226	1226	88.5%
Net-4:	Constrained network 1	2266	1132	94.0%
Net-5:	Constrained network 2	5194	1060	98.4%

- Best results on a large database: Le Cun et al., 1998
  - 60,000 training and 10,000 test examples
  - LeNet-5: a more complex convolutional network
    - 99.2% correct
  - Boosted LeNet-4: boosting with a predecessor of LeNet-5
    - 99.3% correct

# Bibliography

- C. Bishop, Neural Networks for Pattern Recognition. Oxford University Press, 1995.
  - Chapter 4
- T. Hastie, R. Tibshirani, y J. Friedman, The elements of statistical learning. Springer, 2009.
  - Chapter 11
- Unsupervised Feature Learning and Deep Learning Tutorial
  - <http://ufldl.stanford.edu/tutorial/>